

Introduction To

**Machine  
Learning**

**MACHINE LEARNING**

Complete Guide



**GENIAL**  **CODE**

[www.genial-code.com](http://www.genial-code.com)

# SOURCES

- AAI. Machine Learning.  
*<http://www.aaai.org/Pathfinder/html/machine.html>*
- Dietterich, T. (2003). Machine Learning. *Nature Encyclopedia of Cognitive Science*.
- Doyle, P. Machine Learning.  
*<http://www.cs.dartmouth.edu/~brd/Teaching/AI/Lectures/Summaries/learning.html>*
- Dyer, C. (2004). Machine Learning.  
*<http://www.cs.wisc.edu/~dyer/cs540/notes/learning.html>*
- Mitchell, T. (1997). *Machine Learning*.
- Nilsson, N. (2004). Introduction to Machine Learning.  
*<http://robotics.stanford.edu/people/nilsson/mlbook.html>*
- Russell, S. (1997). Machine Learning. *Handbook of Perception and Cognition*, Vol. 14, Chap. 4.
- Russell, S. (2002). *Artificial Intelligence: A Modern Approach*, Chap. 18-20.  
*<http://aima.cs.berkeley.edu>*



# WHAT IS LEARNING?

- *“Learning denotes changes in a system that ... enable a system to do the same task ... more efficiently the next time.”* - Herbert Simon
- *“Learning is constructing or modifying representations of what is being experienced.”* - Ryszard Michalski
- *“Learning is making useful changes in our minds.”* - Marvin Minsky

*“Machine learning refers to a system capable of the autonomous acquisition and integration of knowledge.”*

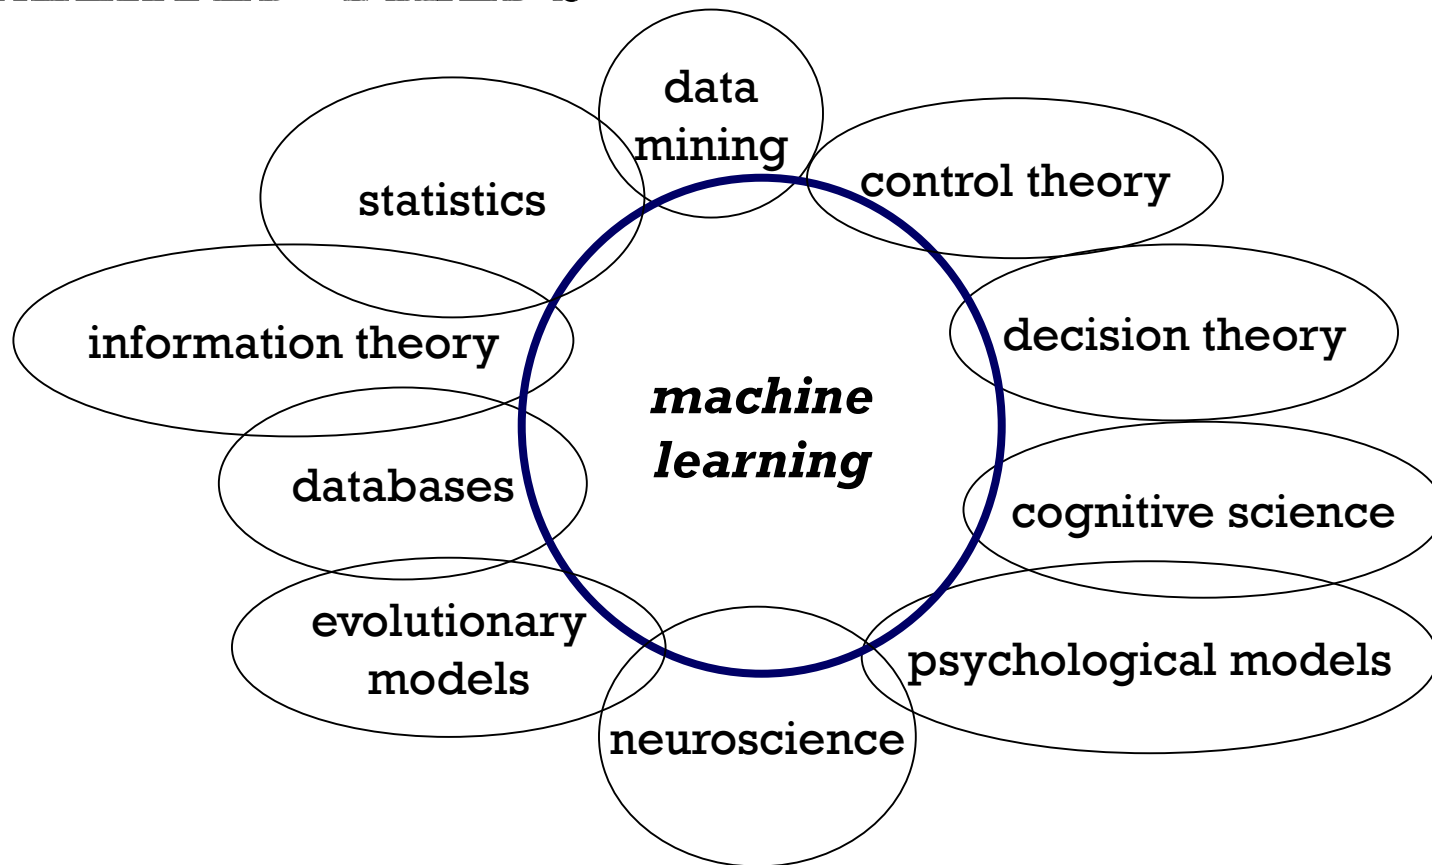


# WHY MACHINE LEARNING?

- No human experts
  - industrial/manufacturing control
  - mass spectrometer analysis, drug design, astronomic discovery
- Black-box human expertise
  - face/handwriting/speech recognition
  - driving a car, flying a plane
- Rapidly changing phenomena
  - credit scoring, financial modeling
  - diagnosis, fraud detection
- Need for customization/personalization
  - personalized news reader
  - movie/book recommendation



# RELATED FIELDS



*Machine learning* is primarily concerned with the accuracy and effectiveness of the *computer system*.

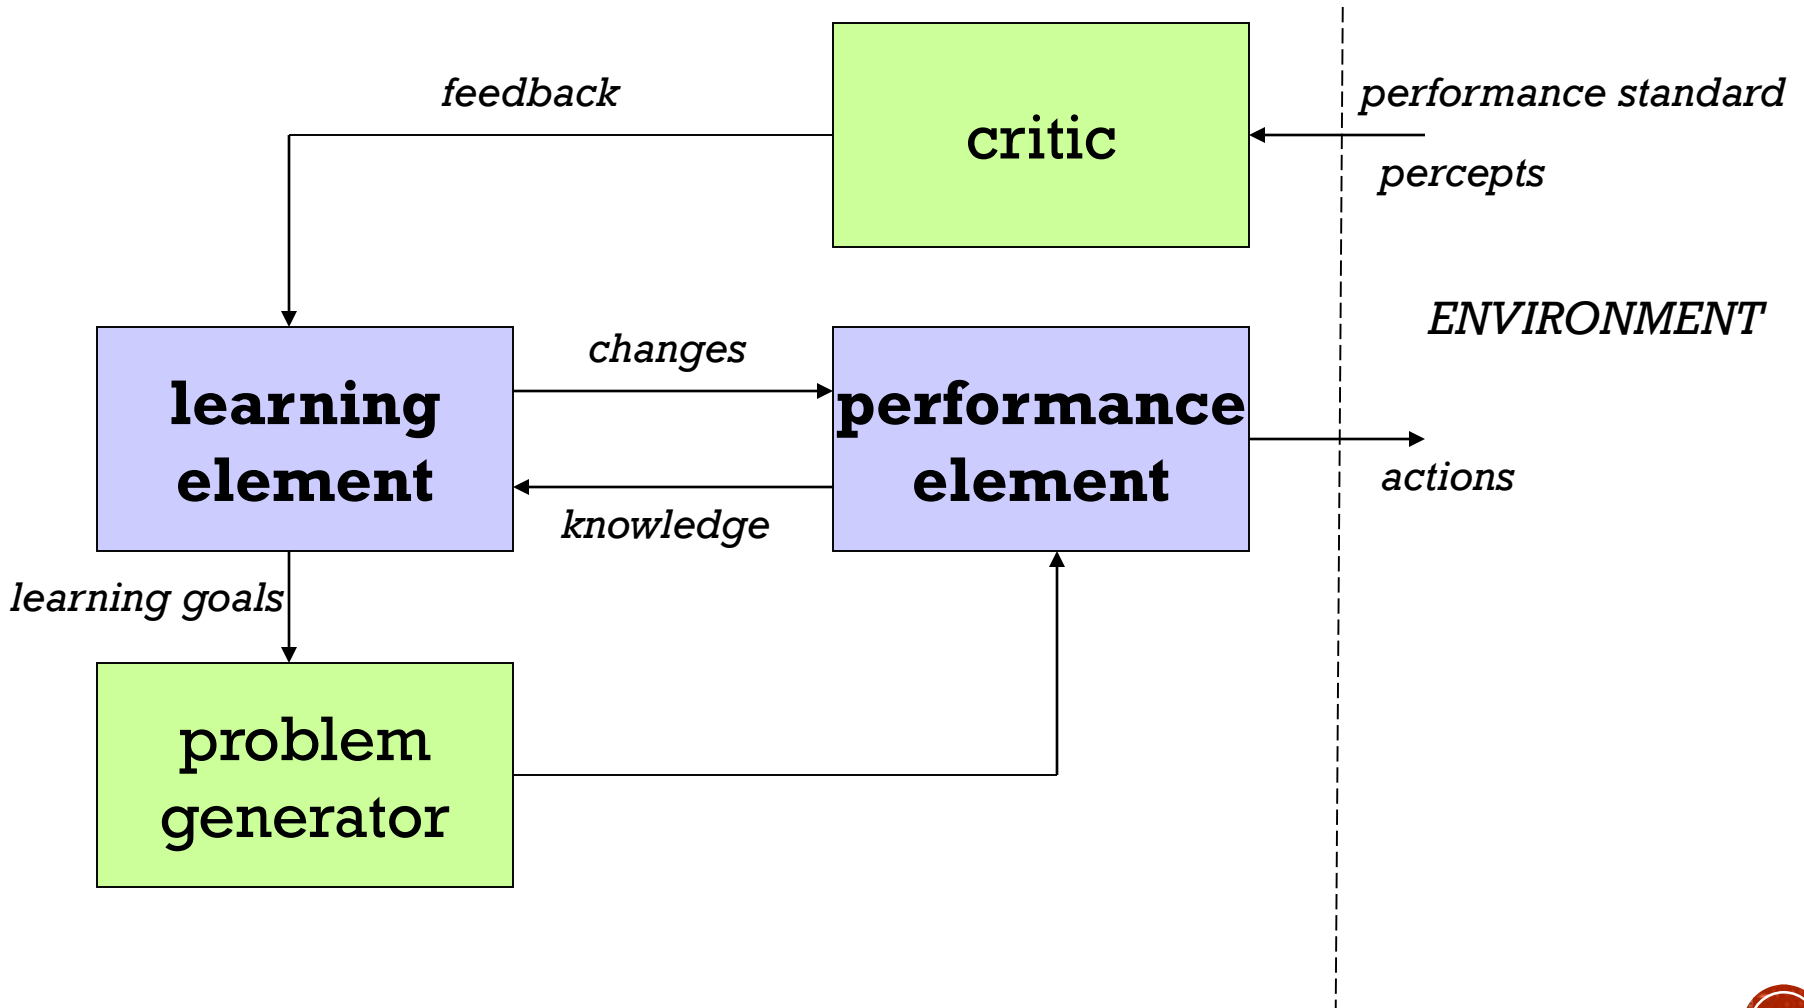


# MACHINE LEARNING PARADIGMS

- **rote learning**
- **learning by being told (advice-taking)**
- **learning from examples (induction)**
- **learning by analogy**
- **speed-up learning**
- **concept learning**
- **clustering**
- **discovery**
- **...**



# ARCHITECTURE OF A LEARNING SYSTEM



# LEARNING ELEMENT

Design affected by:

- *performance element* used
  - e.g., utility-based agent, reactive agent, logical agent
- *functional component* to be learned
  - e.g., classifier, evaluation function, perception-action function,
- *representation* of functional component
  - e.g., weighted linear function, logical theory, HMM
- *feedback* available
  - e.g., correct action, reward, relative preferences





# DIMENSIONS OF LEARNING SYSTEMS

- type of feedback
  - supervised (labeled examples)
  - unsupervised (unlabeled examples)
  - reinforcement (reward)
- representation
  - attribute-based (feature vector)
  - relational (first-order logic)
- use of knowledge
  - empirical (knowledge-free)
  - analytical (knowledge-guided)



# OUTLINE

- Supervised learning
  - empirical learning (knowledge-free)
    - attribute-value representation
    - logical representation
  - analytical learning (knowledge-guided)
- Reinforcement learning
- Unsupervised learning
- *Performance evaluation*
- *Computational learning theory*



# INDUCTIVE (SUPERVISED) LEARNING

Basic Problem: Induce a representation of a function (a systematic relationship between inputs and outputs) from examples.

- **target function**  $f: X \rightarrow Y$
- **example**  $(x, f(x))$
- **hypothesis**  $g: X \rightarrow Y$  such that  $g(x) = f(x)$

$x$  = set of attribute values (**attribute-value representation**)

$x$  = set of logical sentences (*first-order representation*)

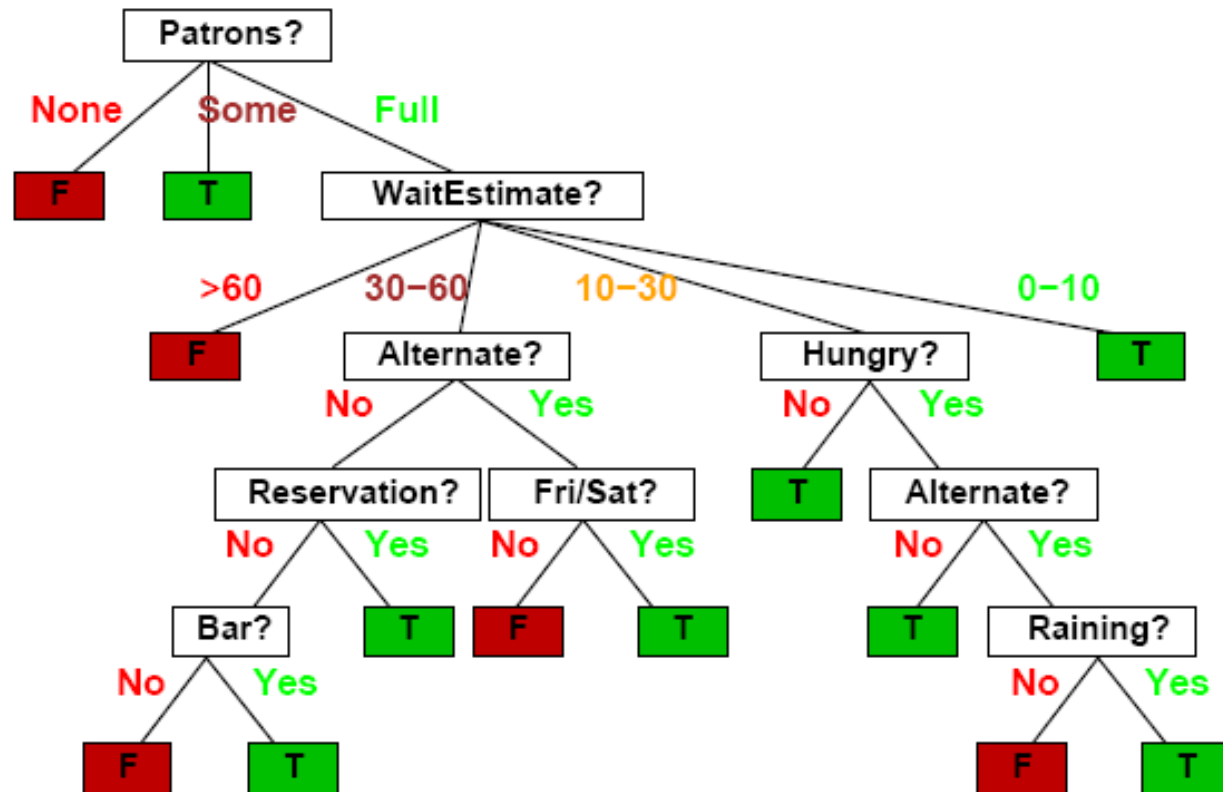
$Y$  = set of discrete labels (**classification**)

$Y = \mathbb{R}$  (*regression*)



# DECISION TREES

*Should I wait at this restaurant?*



# DECISION TREE INDUCTION

(Recursively) partition examples according to the *most important* attribute.



## Key Concepts

- *entropy*
  - impurity of a set of examples (entropy = 0 if perfectly homogeneous)
  - (#bits needed to encode class of an arbitrary example)
- *information gain*
  - expected reduction in entropy caused by partitioning





# DECISION TREE INDUCTION: ATTRIBUTE SELECTION

Intuitively: *A good attribute splits the examples into subsets that are (ideally) all positive or all negative.*




# DECISION TREE INDUCTION: ATTRIBUTE SELECTION

Intuitively: A *good attribute* splits the examples into subsets that are (ideally) *all positive* or *all negative*.




# DECISION TREE INDUCTION: DECISION BOUNDARY







# DECISION TREE INDUCTION: DECISION BOUNDARY



# DECISION TREE INDUCTION: DECISION BOUNDARY




# DECISION TREE INDUCTION: DECISION BOUNDARY




# (ARTIFICIAL) NEURAL NETWORKS

- Motivation: human brain
  - massively parallel ( $10^{11}$  neurons,  $\sim 20$  types)
  - small computational units with simple low-bandwidth communication ( $10^{14}$  synapses, 1-10ms cycle time)
- Realization: neural network
  - *units* ( $\approx$  neurons) connected by *directed weighted links*
  - *activation function* from inputs to output



# NEURAL NETWORKS (*CONTINUED*)



$$\begin{aligned} a_5 &= g(W_{3,5} \cdot a_3 + W_{4,5} \cdot a_4) \\ &= g(W_{3,5} \cdot g(W_{1,3} \cdot a_1 + W_{2,3} \cdot a_2) + W_{4,5} \cdot g(W_{1,4} \cdot a_1 + W_{2,4} \cdot a_2)) \end{aligned}$$

- *neural network = parameterized family of nonlinear functions*
- *types*
  - *feed-forward* (acyclic): single-layer perceptrons, multi-layer networks
  - *recurrent* (cyclic): Hopfield networks, Boltzmann machines

[ *connectionism, parallel distributed processing* ]



# NEURAL NETWORK LEARNING


*Key Idea:* Adjusting the weights changes the function represented by the neural network (*learning = optimization in weight space*).

Iteratively *adjust weights* to reduce error (difference between network output and target output).


- Weight Update
  - *perceptron training rule*
  - *linear programming*
  - *delta rule*
  - *backpropagation*



# NEURAL NETWORK LEARNING: DECISION BOUNDARY



*single-layer perceptron*



*multi-layer network*



# SUPPORT VECTOR MACHINES

*Kernel Trick*: Map data to *higher-dimensional space* where they will be *linearly separable*.

## Learning a Classifier

- optimal linear separator is one that has the *largest margin* between positive examples on one side and negative examples on the other
- = *quadratic programming optimization*





# SUPPORT VECTOR MACHINES (*CONTINUED*)


*Key Concept:* Training data enters optimization problem in the form of *dot products* of pairs of *points*.

- *support vectors*
  - weights associated with data points are *zero* except for those points nearest the separator (i.e., the *support vectors*)
- *kernel function*  $K(\mathbf{x}_i, \mathbf{x}_j)$ 
  - function that can be applied to pairs of points to evaluate dot products in the corresponding (higher-dimensional) feature space  $F$  (*without having to directly compute  $F(\mathbf{x})$  first*)

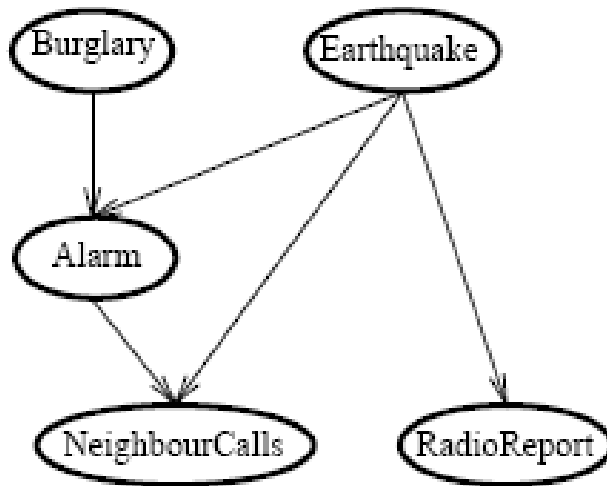
*efficient training and complex functions!*



# SUPPORT VECTOR MACHINES: DECISION BOUNDARY



# BAYESIAN NETWORKS



Network topology reflects direct *causal influence*

Basic Task: Compute probability distribution for unknown variables given observed values of other variables.

	$A B$	$A \neg B$	$\neg A B$	$\neg A \neg B$
$C$	0.9	0.3	0.5	0.1
$\neg C$	0.1	0.7	0.5	0.9

*conditional probability table*  
for NeighbourCalls

[*belief networks, causal networks*]



# BAYESIAN NETWORK LEARNING

## Key Concepts

- nodes (attributes) = random variables
- conditional independence
  - an attribute is conditionally independent of its non-descendants, given its parents
- conditional probability table
  - conditional probability distribution of an attribute given its parents
- Bayes Theorem
  - $P(h | D) = P(D | h)P(h) / P(D)$



# BAYESIAN NETWORK LEARNING

## (*CONTINUED*)

Find *most probable hypothesis* given the data.

*In theory*: Use posterior probabilities to weight hypotheses. (*Bayes optimal classifier*)

*In practice*: Use single, *maximum a posteriori* (most probable) hypothesis.

### Settings

- known structure, fully observable (*parameter learning*)
- unknown structure, fully observable (*structural learning*)
- known structure, hidden variables (*EM algorithm*)
- unknown structure, hidden variables (?)



# NEAREST NEIGHBOR MODELS

*Key Idea:* Properties of an input  $x$  are likely to be *similar* to those of points in the *neighborhood* of  $x$ .



*Basic Idea:* Find ( $k$ ) nearest neighbor(s) of  $x$  and infer target attribute value(s) of  $x$  based on corresponding attribute value(s).

Form of *non-parametric learning* where hypothesis complexity grows with data (learned model  $\approx$  all examples seen so far)

*[instance-based learning, case-based reasoning, analogical reasoning]*



# NEAREST NEIGHBOR MODEL: DECISION BOUNDARY



# LEARNING LOGICAL THEORIES

## Logical Formulation of Supervised Learning

- attribute  $\rightarrow$  unary predicate
- instance  $x \rightarrow$  logical sentence
- positive/negative classifications  $\rightarrow$  sentences  $Q(x_i), \neg Q(x_i)$
- training set  $\rightarrow$  conjunction of all description and classification sentences

*Learning Task:* Find an *equivalent logical expression* for the goal predicate  $Q$  to classify examples correctly.

$$\text{Hypothesis} \wedge \text{Descriptions} \models \text{Classifications}$$





# LEARNING LOGIC THEORIES: EXAMPLE

## Input

- $\text{Father}(\text{Philip}, \text{Charles}), \text{Father}(\text{Philip}, \text{Anne}), \dots$
- $\text{Mother}(\text{Mum}, \text{Margaret}), \text{Mother}(\text{Mum}, \text{Elizabeth}), \dots$
- $\text{Married}(\text{Diana}, \text{Charles}), \text{Married}(\text{Elizabeth}, \text{Philip}), \dots$
- $\text{Male}(\text{Philip}), \text{Female}(\text{Anne}), \dots$
- $\text{Grandparent}(\text{Mum}, \text{Charles}), \text{Grandparent}(\text{Elizabeth}, \text{Beatrice}),$   
 $\neg \text{Grandparent}(\text{Mum}, \text{Harry}), \neg \text{Grandparent}(\text{Spencer}, \text{Pete}), \dots$

## Output

- $\text{Grandparent}(x, y) \Leftrightarrow$   
 $[\exists z \text{Mother}(x, z) \wedge \text{Mother}(z, y)] \vee [\exists z \text{Mother}(x, z) \wedge \text{Father}(z, y)] \vee$   
 $[\exists z \text{Father}(x, z) \wedge \text{Mother}(z, y)] \vee [\exists z \text{Father}(x, z) \wedge \text{Father}(z, y)]$



# LEARNING LOGIC THEORIES

## Key Concepts


- *specialization*
  - triggered by false positives (*goal: exclude negative examples*)
  - achieved by adding conditions, dropping disjuncts
- *generalization*
  - triggered by false negatives (*goal: include positive examples*)
  - achieved by dropping conditions, adding disjuncts

## Learning


- *current-best-hypothesis*: incrementally improve single hypothesis (e.g., *sequential covering*)
- *least-commitment search*: maintain *all* hypotheses consistent with examples seen so far (e.g., *version space*)




# LEARNING LOGIC THEORIES: DECISION BOUNDARY




# LEARNING LOGIC THEORIES: DECISION BOUNDARY




# LEARNING LOGIC THEORIES: DECISION BOUNDARY



# LEARNING LOGIC THEORIES: DECISION BOUNDARY



# LEARNING LOGIC THEORIES: DECISION BOUNDARY



# ANALYTICAL LEARNING

## Prior Knowledge in Learning

*Recall:*

$\text{Grandparent}(x,y) \Leftrightarrow$

$[\exists z \text{ Mother}(x,z) \wedge \text{Mother}(z,y)] \vee [\exists z \text{ Mother}(x,z) \wedge \text{Father}(z,y)] \vee$   
 $[\exists z \text{ Father}(x,z) \wedge \text{Mother}(z,y)] \vee [\exists z \text{ Father}(x,z) \wedge \text{Father}(z,y)]$

- Suppose initial theory also included:
  - $\text{Parent}(x,y) \Leftrightarrow [\text{Mother}(x,y) \vee \text{Father}(x,y)]$
- Final Hypothesis:
  - $\text{Grandparent}(x,y) \Leftrightarrow [\exists z \text{ Parent}(x,z) \wedge \text{Parent}(z,y)]$

***Background knowledge can dramatically reduce the size of the hypothesis (greatly simplifying the learning problem).***





# EXPLANATION-BASED LEARNING

Amazed crowd of cavemen observe Zog roasting a lizard on the end of a pointed stick ("Look what Zog do!") and thereafter abandon roasting with their bare hands.

*Basic Idea:* Generalize by *explaining* observed instance.

- form of *speedup learning*
  - doesn't learn anything factually new from the observation
  - instead converts first-principles theories into *useful* special-purpose knowledge
- utility problem
  - cost of determining if learned knowledge is applicable may outweigh benefits from its application



# RELEVANCE-BASED LEARNING

Mary travels to Brazil and meets her first Brazilian (Fernando), who speaks Portuguese. She concludes that all Brazilians speak Portuguese but not that all Brazilians are named Fernando.

*Basic Idea:* Use knowledge of what is *relevant* to infer new properties about a new instance.

- form of *deductive learning*
  - learns a new general rule that explains observations
  - does not create knowledge outside logical content of prior knowledge and observations



# KNOWLEDGE-BASED INDUCTIVE LEARNING

Medical student observes consulting session between doctor and patient at the end of which the doctor prescribes a particular medication. Student concludes that the medication is effective treatment for a particular type of infection.

Basic Idea: Use prior knowledge to *guide hypothesis generation*.

- benefits in inductive logic programming
  - only hypotheses consistent with prior knowledge and observations are considered
  - prior knowledge supports smaller (simpler) hypotheses



# REINFORCEMENT LEARNING

*k*-armed bandit problem:

*Agent is in a room with  $k$  gambling machines (one-armed bandits). When an arm is pulled, the machine pays off 1 or 0, according to some unknown probability distribution. Given a fixed number of pulls, what is the agent's (optimal) strategy?*

**Basic Task:** Find a policy  $\pi$ , mapping states to actions, that maximizes (long-term) reward.

Model (*Markov Decision Process*)

- set of states  $S$
- set of actions  $A$
- reward function  $R : S \times A \rightarrow \mathbb{R}$
- state transition function  $T : S \times A \rightarrow \Pi(S)$ 
  - $T(s,a,s')$  = probability of reaching  $s'$  when  $a$  is executed in  $s$



# REINFORCEMENT LEARNING (*CONTINUED*)

- Settings
  - fully vs. partially observable environment
  - deterministic vs. stochastic environment
  - model-based vs. model-free
  - rewards in goal state only or in any state

*value of a state*: expected *infinite discounted sum of reward* the agent will gain if it starts from that state and *executes the optimal policy*

Solving MDP when the model is known

- *value iteration*: find optimal value function (derive optimal policy)
- *policy iteration*: find optimal policy directly (derive value function)



# REINFORCEMENT LEARNING (*CONTINUED*)

Reinforcement learning is concerned with finding an optimal policy for an MDP when the *model* (transition, reward) is *unknown*.

*exploration/exploitation tradeoff*

model-free reinforcement learning

- learn a controller without learning a model first
- e.g., *adaptive heuristic critic* ( $TD(\lambda)$ ), *Q-learning*

model-based reinforcement learning

- learn a model first
- e.g., *Dyna*, *prioritized sweeping*, *RTDP*



# UNSUPERVISED LEARNING

*Learn patterns from (unlabeled) data.*

## Approaches

- clustering (similarity-based)
- density estimation (e.g., EM algorithm)

## Performance Tasks

- understanding and visualization
- anomaly detection
- information retrieval
- data compression




# PERFORMANCE EVALUATION

- Randomly split examples into *training set*  $U$  and *test set*  $V$ .
- Use training set to learn a hypothesis  $H$ .
- Measure % of  $V$  correctly classified by  $H$ .
- Repeat for different random splits and average results.






# PERFORMANCE EVALUATION: LEARNING CURVES



# PERFORMANCE EVALUATION: ROC CURVES




# PERFORMANCE EVALUATION: ACCURACY/COVERAGE



# TRIPLE TRADEOFF IN EMPIRICAL LEARNING

- size/complexity of learned classifier
- amount of training data
- generalization accuracy

*bias-variance tradeoff*



# COMPUTATIONAL LEARNING THEORY

*probably approximately correct (PAC) learning*

With probability  $\geq 1 - \delta$ , error will be  $\leq \epsilon$ .

*Basic principle:* Any hypothesis that is seriously wrong will almost certainly be found out with high probability after a small number of examples.

*Key Concepts*

$$m \geq \frac{1}{\epsilon} \left( \ln \frac{1}{\delta} + \ln |H| \right)$$

- examples drawn from same distribution (*stationarity assumption*)
- *sample complexity* is a function of confidence, error, and size of hypothesis space



# CURRENT MACHINE LEARNING RESEARCH

- Representation
  - data sequences
  - spatial/temporal data
  - probabilistic relational models
  - ...
- Approaches
  - ensemble methods
  - cost-sensitive learning
  - active learning
  - semi-supervised learning
  - collective classification
  - ...

