


# Reference Python Notebook for Important ML Algorithms for the Beginners

<https://ashutoshtrpathi.com/>

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 180212 entries, 0 to 180211
Data columns (total 11 columns):
patient_id      180212 non-null int64
Age             180212 non-null int64
Gender          180212 non-null object
Prescription_period 180212 non-null int64
Diabetes        180212 non-null int64
Alcoholism      180212 non-null int64
HyperTension    180212 non-null int64
Smokes         180212 non-null int64
Tuberculosis    180212 non-null int64
Sms_Reminder    180212 non-null int64
Adherence       180212 non-null object
dtypes: int64(9), object(2)
memory usage: 15.1+ MB
```



```
learning rate: 0.75
Accuracy score (training): 0.895
Accuracy score (testing): 0.893
#### Classification Report ####
              precision    recall  f1-score   support

0               0.95         0.90         0.92       25163
1               0.79         0.88         0.83       10880

 accuracy         0.87         0.89         0.89       36043
 macro avg        0.87         0.89         0.88       36043
 weighted avg     0.90         0.89         0.89       36043

#### Confusion Matrix ####
[[22594 2569]
 [ 1284  9596]]

# ROC curve and Area-under-Curve (AUC)
y_scores_gb = gb.decision_function(X_test)
fpr_gb, tpr_gb, _ = roc_curve(y_test, y_scores_gb)
roc_auc_gb = auc(fpr_gb, tpr_gb)

print("Area under ROC curve = {:.2f}".format(roc_auc_gb))
```

Data set

Download

HTML File

Download

Python Notebook

Download

Download link: <https://ashutoshtrpathi.com/reference-python-notebook-for-important-ml-algorithms-for-the-beginners/>

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
```

## Exploratory Data Analysis

```
In [24]: data = pd.read_csv('Training Data.csv')
```

```
In [3]: data.head() #produce top 10 rows
```

```
Out[3]:
```

	patient_id	Age	Gender	Prescription_period	Diabetes	Alcoholism	HyperTension	Smokes	Tuberculosis	Sms_Reminder	Adherence
0	1	19	M	7	0	0	0	0	0	0	No
1	2	24	F	59	0	0	0	0	0	0	No
2	3	4	F	43	0	0	0	0	0	0	No
3	4	38	M	66	0	0	0	0	0	1	No
4	5	46	F	98	0	0	0	0	0	1	No

```
In [4]: data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 180212 entries, 0 to 180211
Data columns (total 11 columns):
patient_id      180212 non-null int64
Age             180212 non-null int64
Gender          180212 non-null object
Prescription_period 180212 non-null int64
Diabetes        180212 non-null int64
Alcoholism      180212 non-null int64
HyperTension    180212 non-null int64
Smokes         180212 non-null int64
Tuberculosis    180212 non-null int64
Sms_Reminder    180212 non-null int64
Adherence       180212 non-null object
dtypes: int64(9), object(2)
memory usage: 15.1+ MB
```

in the above info, we see none of the column has missing values as all have a total of 180212 entries. So will not go for missing value treatment

```
In [5]: data.describe()
```

Out[5]:

	patient_id	Age	Prescription_period	Diabetes	Alcoholism	HyperTension	Smokes	Tuberculosis	Sms_Re
count	180212.000000	180212.000000	180212.000000	180212.000000	180212.000000	180212.000000	180212.000000	180212.000000	180212.000000
mean	90106.500000	37.795363	54.668485	0.078524	0.025043	0.216512	0.052566	0.000338	0.57396
std	52022.867693	22.852072	35.752491	0.268995	0.156255	0.411868	0.223166	0.018395	0.49982
min	1.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00000
25%	45053.750000	19.000000	22.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00000
50%	90106.500000	38.000000	51.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.00000
75%	135159.250000	56.000000	86.000000	0.000000	0.000000	0.000000	0.000000	0.000000	1.00000
max	180212.000000	113.000000	120.000000	1.000000	1.000000	1.000000	1.000000	1.000000	2.00000

```
In [6]: print(data.groupby('Adherence').size())
```

```
Adherence
No      125822
Yes     54390
dtype: int64
```

**Target class is imbalance. However I will not go for Sampling techniques, as performance criteria is not the accuracy. In Medicals, We mostly focus on Precision and Recall.**

```
In [7]: print(data.shape)
```

```
(180212, 11)
```

```
In [8]: data.isnull().count() # no null values
```

Out[8]:

patient_id	180212
Age	180212
Gender	180212
Prescription_period	180212
Diabetes	180212
Alcoholism	180212
HyperTension	180212
Smokes	180212
Tuberculosis	180212
Sms_Reminder	180212
Adherence	180212

dtype: int64

```
In [9]: data.isna().count() # no missing values
```

Out[9]:

patient_id	180212
Age	180212
Gender	180212
Prescription_period	180212
Diabetes	180212
Alcoholism	180212
HyperTension	180212
Smokes	180212
Tuberculosis	180212
Sms_Reminder	180212
Adherence	180212

dtype: int64

## Data Visualization

```
In [31]: !pip install plotly
```


```
In [10]: import plotly
import plotly.offline as pyoff
import plotly.figure_factory as ff
from plotly.offline import init_notebook_mode, iplot, plot
import plotly.graph_objs as go
```

```
In [11]: temp = data.Adherence.value_counts()
trace = go.Bar(x=temp.index,
               y= np.round(temp.astype(float)/temp.values.sum(),2),
               text = np.round(temp.astype(float)/temp.values.sum(),2),
               textposition = 'auto',
               name = 'Adherence')

data1 = [trace]
layout = go.Layout(
    autosize=False,
    width=600,
    height=400, title = "Adherence Distribution"
)

fig = go.Figure(data=data1, layout=layout)
iplot(fig)
del temp
```


Adherence Distribution



```
In [12]: print(data.groupby('Gender').size())
print('-----')
sns.factorplot('Gender', data=data, kind='count')
```

```
Gender
F    120438
M     59774
dtype: int64
-----
```

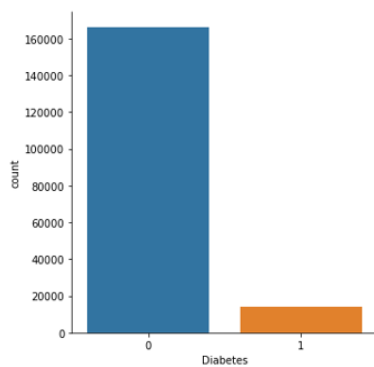
```
Out[12]: <seaborn.axisgrid.FacetGrid at 0x168788b66d8>
```



```
In [13]: print(data.groupby('Diabetes').size())
print('-----')
sns.factorplot('Diabetes', data=data, kind='count')
```

```
Diabetes
0    166061
1     14151
dtype: int64
-----
```

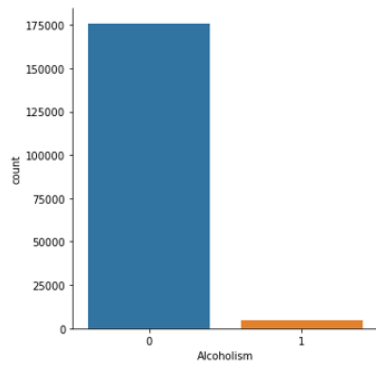
```
Out[13]: <seaborn.axisgrid.FacetGrid at 0x1680e600b70>
```



```
In [14]: print(data.groupby('Alcoholism').size())
print('-----')
sns.factorplot('Alcoholism', data=data, kind='count')
```

```
Alcoholism
0    175699
1     4513
dtype: int64
-----
```

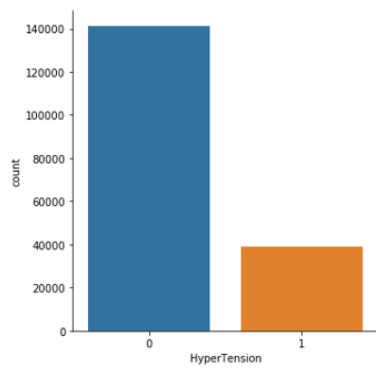
Out[14]: <seaborn.axisgrid.FacetGrid at 0x1680c770cc0>



```
In [15]: print(data.groupby('HyperTension').size())
sns.factorplot('HyperTension', data=data, kind='count')
```

```
HyperTension
0    141194
1     39018
dtype: int64
```


Out[15]: <seaborn.axisgrid.FacetGrid at 0x1680e6bf8d0>



```
In [16]: print(data.groupby('Smokes').size())
sns.factorplot('Smokes', data=data, kind='count')
```

```
Smokes
0    170739
1     9473
dtype: int64
```


Out[16]: <seaborn.axisgrid.FacetGrid at 0x1680e72dfd0>



```
In [17]: print(data.groupby('Tuberculosis').size())
sns.factorplot('Tuberculosis', data=data, kind='count')
```

```
Tuberculosis
0    180151
1         61
dtype: int64
```

Out[17]: <seaborn.axisgrid.FacetGrid at 0x1680c770d30>



## Multivariate Analysis and Visualization


```
In [25]: Gender_Adherence = data.groupby(['Gender','Adherence']).size().to_frame()
Gender_Adherence = Gender_Adherence.reset_index()
Gender_Adherence.columns = ['Gender','Adherence','Count']
Gender_Adherence
```

Out[25]:

	Gender	Adherence	Count
0	F	No	84487
1	F	Yes	35951
2	M	No	41335
3	M	Yes	18439

```
In [26]: trace1 = go.Bar(x = Gender_Adherence.Gender[Gender_Adherence.Adherence=='Yes'],
                        y = Gender_Adherence.Count[Gender_Adherence.Adherence=='Yes'],
                        text = Gender_Adherence.Count[Gender_Adherence.Adherence=='Yes'],
                        textposition = 'auto',
                        name = 'Yes')
trace2 = go.Bar(x = Gender_Adherence.Gender[Gender_Adherence.Adherence=='No'],
                y = Gender_Adherence.Count[Gender_Adherence.Adherence=='No'],
                text = Gender_Adherence.Count[Gender_Adherence.Adherence=='No'],
                textposition = 'auto',
                name = 'No')
tempdata = [trace1,trace2]
layout = go.Layout(width = 800,
                  height = 600,title = 'Gender and Adherence')
fig = go.Figure(data=tempdata, layout=layout)
iplot(fig)
```


Gender and Adherence



```
In [27]: print(data.groupby(['Tuberculosis','Adherence']).size())
sns.factorplot('Tuberculosis',data=data,hue='Adherence',kind='count')
```

```
Tuberculosis  Adherence
0             No      125783
              Yes       54368
1             No         39
              Yes         22
dtype: int64
```


Out[27]: <seaborn.axisgrid.FacetGrid at 0x168106f00b8>



```
In [28]: print(data.groupby(['Smokes', 'Adherence']).size())
sns.factorplot('Smokes', data=data, hue='Adherence', kind='count')
```

```
Smokes  Adherence
0       No         119651
       Yes         51088
1       No          6171
       Yes          3302
dtype: int64
```


Out[28]: <seaborn.axisgrid.FacetGrid at 0x1680e6eeb38>



```
In [29]: print(data.groupby(['HyperTension', 'Adherence']).size())
sns.catplot('HyperTension', data=data, hue='Adherence', kind='count')
```

```
HyperTension  Adherence
0            No         96491
             Yes         44703
1            No         29331
             Yes          9687
dtype: int64
```


Out[29]: <seaborn.axisgrid.FacetGrid at 0x1680dbeef98>



```
In [30]: print(data.groupby(['Alcoholism', 'Adherence']).size())
sns.factorplot('Alcoholism', data=data, hue='Adherence', kind='count')
```

```
Alcoholism  Adherence
0          No        122994
           Yes         52705
1          No         2828
           Yes         1685
dtype: int64
```


Out[30]: <seaborn.axisgrid.FacetGrid at 0x16810cdbf98>



```
In [31]: print(data.groupby(['Diabetes', 'Adherence']).size())
sns.factorplot('Diabetes', data=data, hue='Adherence', kind='count')
```

```
Diabetes  Adherence
0         No         115252
         Yes         50809
1         No         10570
         Yes          3581
dtype: int64
```

Out[31]: <seaborn.axisgrid.FacetGrid at 0x16810c5cfd0>



## Data pre-processing

Convert the Adherence and Gender column to integer 0 and 1.

Adherence Yes = 1 and No = 0, Gender M = 1, F = 0

Then, I set the patient\_id column to be the index of the dataframe.

After all, the patient\_id column will not be used as all values are unique

```
In [32]: data['Adherence'] = data['Adherence'].apply(lambda x: 1 if x == 'Yes' else 0)
data['Gender'] = data['Gender'].apply(lambda x: 1 if x == 'M' else 0)
data = data.set_index('patient_id')
```

```
In [33]: data.head()
```

```
Out[33]:
```

	Age	Gender	Prescription_period	Diabetes	Alcoholism	HyperTension	Smokes	Tuberculosis	Sms_Reminder	Adherence
patient_id										
1	19	1	7	0	0	0	0	0	0	0
2	24	0	59	0	0	0	0	0	0	0
3	4	0	43	0	0	0	0	0	0	0
4	38	1	66	0	0	0	0	0	1	0
5	46	0	98	0	0	0	0	0	1	0

```
In [34]: data.groupby(['Diabetes', 'Adherence']).size()
```

```
Out[34]: Diabetes  Adherence
0                0      115252
           1        50809
1                0      10570
           1         3581
dtype: int64
```

```
In [35]: data.groupby(['Diabetes', 'Alcoholism', 'Adherence']).size()
```

```
Out[35]: Diabetes  Alcoholism  Adherence
0                0              0      112792
           1              1      49277
           0              0      2460
           1              1      1532
1                0              0      10202
           1              1      3428
           0              0       368
           1              1       153
dtype: int64
```

```
In [36]: data.groupby('Adherence').size()
```

```
Out[36]: Adherence
0      125822
1       54390
dtype: int64
```

## Train Test Split

```
In [37]: from sklearn.model_selection import train_test_split
```

```
In [38]: X_train, X_test, Y_train, Y_test = train_test_split(data.drop('Adherence', axis = 1), data['Adherence'], test_size=0.20, random_state=101)
```

```
In [39]: X_train.head()
```

```
Out[39]:
```

	Age	Gender	Prescription_period	Diabetes	Alcoholism	HyperTension	Smokes	Tuberculosis	Sms_Reminder
patient_id									
116942	69	1	1	1	0	1	0	0	1
52438	69	0	78	0	0	0	0	0	0
110751	80	0	55	0	0	1	0	0	0
32426	41	0	99	0	0	0	0	0	1
81027	55	0	10	0	0	1	0	0	0

```
In [40]: X_test.head()
```

```
Out[40]:
```

	Age	Gender	Prescription_period	Diabetes	Alcoholism	HyperTension	Smokes	Tuberculosis	Sms_Reminder
patient_id									
161759	18	0	93	0	0	0	0	0	1
131777	30	0	118	0	0	1	0	0	1
22965	55	0	103	0	0	1	0	0	1
1468	41	0	69	0	0	0	0	0	0
170654	1	0	71	0	0	0	0	0	0

```
In [41]: Y_train.head()
```

```
Out[41]: patient_id
116942    1
52438     0
110751    0
32426     0
81027     1
Name: Adherence, dtype: int64
```



```
In [42]: Y_test.head()
Out[42]: patient_id
161759    0
131777    0
22965     0
1468      0
170654    0
Name: Adherence, dtype: int64
```

## Model Building | Try different Models and check Performance Metrics

### Logistic Regression Classifier

```
In [43]: from sklearn.linear_model import LogisticRegression
In [44]: logClassifier = LogisticRegression()
In [45]: logClassifier.fit(X_train, Y_train)
Out[45]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
intercept_scaling=1, l1_ratio=None, max_iter=100,
multi_class='auto', n_jobs=None, penalty='l2',
random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
warm_start=False)
In [46]: predictions = logClassifier.predict(X_test)
In [47]: from sklearn.metrics import classification_report, confusion_matrix, roc_curve, auc
In [48]: print('##### Classification Report of Logistic Regression #####')
print(classification_report(Y_test,predictions))
print('##### Confusion Matrix of Logistic Regression #####')
print(confusion_matrix(Y_test,predictions))

##### Classification Report of Logistic Regression #####
              precision    recall  f1-score   support

      0       0.92      0.91      0.91      25163
      1       0.79      0.82      0.80      10880

   accuracy      0.86      0.86      0.88      36043
  macro avg      0.86      0.86      0.86      36043
weighted avg      0.88      0.88      0.88      36043

##### Confusion Matrix of Logistic Regression #####
[[22808  2355]
 [ 1986  8894]]
```

### KNN Classifier

```
In [49]: from sklearn.neighbors import KNeighborsClassifier
knn_3 = KNeighborsClassifier(n_neighbors=3)
In [50]: knn_3.fit(X_train, Y_train)
knn_pred = knn_3.predict(X_test)
In [51]: print('##### Classification Report of KNN Classifier #####')
print()
print(classification_report(Y_test,knn_pred))
print('##### Confusion Matrix of KNN Classifier #####')
print()
print(confusion_matrix(Y_test,knn_pred))

##### Classification Report of KNN Classifier #####
              precision    recall  f1-score   support

      0       0.91      0.90      0.90      25163
      1       0.78      0.79      0.78      10880

   accuracy      0.84      0.84      0.87      36043
  macro avg      0.84      0.84      0.84      36043
weighted avg      0.87      0.87      0.87      36043

##### Confusion Matrix of KNN Classifier #####
[[22694  2469]
 [ 2316  8564]]
```

### Decision Tree Classifier

```
In [52]: from sklearn.tree import DecisionTreeClassifier
In [53]: dt = DecisionTreeClassifier()
In [54]: dt.fit(X_train, Y_train)
dt_pred = dt.predict(X_test)
In [80]: print('##### Classification Report of Decision Tree Classifier #####')
print()
print(classification_report(Y_test,dt_pred))
print('##### Confusion Matrix of Decision Tree Classifier #####')
print()
print(confusion_matrix(Y_test,dt_pred))
```

##### Classification Report of Decision Tree Classifier #####

	precision	recall	f1-score	support
0	0.89	0.90	0.90	25163
1	0.76	0.75	0.76	10880
accuracy			0.85	36043
macro avg	0.83	0.83	0.83	36043
weighted avg	0.85	0.85	0.85	36043

##### Confusion Matrix of Decision Tree Classifier #####

```
[[22568 2595]
 [ 2684 8196]]
```

## Ensemble Learning | Bagging with Random Feature Subspaces that is Random Forest Classifier

```
In [55]: from sklearn.ensemble import RandomForestClassifier
```

```
In [56]: rf = RandomForestClassifier()
```

```
In [57]: rf.fit(X_train, Y_train)
rf_pred = rf.predict(X_test)
```

```
In [58]: print('##### Classification Report of Random Forest Classifier #####')
print()
print(classification_report(Y_test, rf_pred))
print('##### Confusion Matrix of Random Forest Classifier #####')
print()
print(confusion_matrix(Y_test, rf_pred))
```

##### Classification Report of Random Forest Classifier #####

	precision	recall	f1-score	support
0	0.92	0.90	0.91	25163
1	0.77	0.81	0.79	10880
accuracy			0.87	36043
macro avg	0.84	0.85	0.85	36043
weighted avg	0.87	0.87	0.87	36043

##### Confusion Matrix of Random Forest Classifier #####

```
[[22534 2629]
 [ 2075 8805]]
```

## Ensemble Learning | Improving Weak Learners with Boosting

### Adaboost Classifier

```
In [59]: from sklearn.ensemble import AdaBoostClassifier
```

```
In [60]: classifier = AdaBoostClassifier(
DecisionTreeClassifier(max_depth=1),
n_estimators=200
)
```

```
In [61]: classifier.fit(X_train, Y_train)
```

```
Out[61]: AdaBoostClassifier(algorithm='SAMME.R',
base_estimator=DecisionTreeClassifier(ccp_alpha=0.0,
class_weight=None,
criterion='gini',
max_depth=1,
max_features=None,
max_leaf_nodes=None,
min_impurity_decrease=0.0,
min_impurity_split=None,
min_samples_leaf=1,
min_samples_split=2,
min_weight_fraction_leaf=0.0,
presort='deprecated',
random_state=None,
splitter='best'),
learning_rate=1.0, n_estimators=200, random_state=None)
```

```
In [62]: adaboost_pred = classifier.predict(X_test)
```

```
In [63]: print('##### Classification Report of Adaboost Classifier #####')
print()
print(classification_report(Y_test, adaboost_pred))
print('##### Confusion Matrix of Adaboost Classifier #####')
print()
print(confusion_matrix(Y_test, adaboost_pred))
```

##### Classification Report of Adaboost Classifier #####

	precision	recall	f1-score	support
0	0.95	0.89	0.92	25163
1	0.79	0.89	0.83	10880
accuracy			0.89	36043
macro avg	0.87	0.89	0.88	36043
weighted avg	0.90	0.89	0.89	36043

##### Confusion Matrix of Adaboost Classifier #####

```
[[22517 2646]
 [ 1217 9663]]
```

## Gradient Boosting Classifier | Multiple Iterations with different Learning Rate

```
In [64]: from sklearn.ensemble import GradientBoostingClassifier
```

```
In [65]: learning_rates = [0.05, 0.1, 0.25, 0.5, 0.75, 1]
for learning_rate in learning_rates:
    gb = GradientBoostingClassifier(n_estimators=20, learning_rate = learning_rate, max_features=2, max_depth = 2, random_state
    = 0)
    gb.fit(X_train, Y_train)
    print("Learning rate: ", learning_rate)
    print("Accuracy score (training): {0:.3f}".format(gb.score(X_train, Y_train)))
    print("Accuracy score (testing): {0:.3f}".format(gb.score(X_test, Y_test)))
    gb_pred = gb.predict(X_test)
    print('Classification Report')
    print(classification_report(Y_test,gb_pred))
    print('Confusion Matrix')
    print(confusion_matrix(Y_test,gb_pred))
    print('-----')
```

```
Learning rate: 0.05
Accuracy score (training): 0.701
Accuracy score (testing): 0.702
Classification Report
      precision    recall  f1-score   support

     0       0.70      1.00      0.82      25163
     1       0.88      0.01      0.03      10880

 accuracy      0.79      0.51      0.70      36043
 macro avg      0.79      0.51      0.43      36043
weighted avg      0.75      0.70      0.58      36043
```

```
Confusion Matrix
[[25142   21]
 [10733  147]]
```

```
-----
Learning rate: 0.1
Accuracy score (training): 0.895
Accuracy score (testing): 0.893
Classification Report
      precision    recall  f1-score   support

     0       0.95      0.90      0.92      25163
     1       0.79      0.88      0.83      10880

 accuracy      0.87      0.89      0.89      36043
 macro avg      0.87      0.89      0.88      36043
weighted avg      0.90      0.89      0.89      36043
```

```
Confusion Matrix
[[22552 2611]
 [ 1257 9623]]
```

```
-----
Learning rate: 0.25
Accuracy score (training): 0.895
Accuracy score (testing): 0.893
Classification Report
      precision    recall  f1-score   support

     0       0.95      0.89      0.92      25163
     1       0.78      0.89      0.83      10880

 accuracy      0.87      0.89      0.89      36043
 macro avg      0.87      0.89      0.88      36043
weighted avg      0.90      0.89      0.89      36043
```

```
Confusion Matrix
[[22490 2673]
 [ 1201 9679]]
```

```
-----
Learning rate: 0.5
Accuracy score (training): 0.895
Accuracy score (testing): 0.893
Classification Report
      precision    recall  f1-score   support

     0       0.95      0.90      0.92      25163
     1       0.79      0.88      0.83      10880

 accuracy      0.87      0.89      0.89      36043
 macro avg      0.87      0.89      0.88      36043
weighted avg      0.90      0.89      0.89      36043
```

```
Confusion Matrix
[[22576 2587]
 [ 1273 9607]]
```

```
-----
Learning rate: 0.75
Accuracy score (training): 0.895
Accuracy score (testing): 0.893
Classification Report
      precision    recall  f1-score   support

     0       0.95      0.90      0.92      25163
     1       0.79      0.88      0.83      10880

 accuracy      0.87      0.89      0.89      36043
 macro avg      0.87      0.89      0.88      36043
weighted avg      0.90      0.89      0.89      36043
```

```
Confusion Matrix
[[22594 2569]
 [ 1284 9596]]
```

```

Learning rate: 1
Accuracy score (training): 0.895
Accuracy score (testing): 0.893
Classification Report
      precision    recall  f1-score   support

      0       0.95      0.89      0.92      25163
      1       0.78      0.89      0.83      10880

   accuracy       0.89      0.89      0.89      36043
  macro avg       0.87      0.89      0.88      36043
 weighted avg       0.90      0.89      0.89      36043

Confusion Matrix
[[22507  2656]
 [ 1206  9674]]
-----

```

## Final Model with Gradient Boosting Classifier with Hyperparameter tuned with Learning Rate = 0.75

```

In [66]: learning_rate = 0.75
gb = GradientBoostingClassifier(n_estimators=20, learning_rate = learning_rate, max_features=2, max_depth = 2, random_state = 0)
gb.fit(X_train, Y_train)
print("Learning rate: ", learning_rate)
print("Accuracy score (training): {0:.3f}".format(gb.score(X_train, Y_train)))
print("Accuracy score (testing): {0:.3f}".format(gb.score(X_test, Y_test)))
gb_pred = gb.predict(X_test)
print('##### Classification Report #####')
print(classification_report(Y_test,gb_pred))
print('##### Confusion Matrix #####')
print(confusion_matrix(Y_test,gb_pred))

Learning rate: 0.75
Accuracy score (training): 0.895
Accuracy score (testing): 0.893
##### Classification Report #####
      precision    recall  f1-score   support

      0       0.95      0.90      0.92      25163
      1       0.79      0.88      0.83      10880

   accuracy       0.89      0.89      0.89      36043
  macro avg       0.87      0.89      0.88      36043
 weighted avg       0.90      0.89      0.89      36043

##### Confusion Matrix #####
[[22594  2569]
 [ 1284  9596]]

```

```

In [67]: # ROC curve and Area-Under-Curve (AUC)
y_scores_gb = gb.decision_function(X_test)
fpr_gb, tpr_gb, _ = roc_curve(Y_test, y_scores_gb)
roc_auc_gb = auc(fpr_gb, tpr_gb)

print("Area under ROC curve = {:.2f}".format(roc_auc_gb))

Area under ROC curve = 0.91

```

## Prediction on the Test Data Set Provided | Generating final Result.csv in given format

```
In [68]: test_data = pd.read_csv('Test Data.csv')
```

```
In [69]: test_data['Gender'] = test_data['Gender'].apply(lambda x: 1 if x == 'M' else 0)
```

```
In [70]: test_data = test_data.set_index('patient_id')
```

```
In [71]: test_data.head()
```

```
Out[71]:
```

	Age	Gender	Prescription_period	Diabetes	Alcoholism	HyperTension	Smokes	Tuberculosis	Sms_Reminder
patient_id									
1	5	1	28	0	0	0	0	0	1
2	62	0	9	1	0	1	0	0	0
3	4	0	73	0	0	0	0	0	1
4	33	1	117	0	0	0	0	0	0
5	38	1	8	0	0	0	0	0	1

```
In [72]: predictions_test = gb.predict(test_data)
```

```
In [73]: predictedProbabilityScoresForEachClass = gb.predict_proba(test_data)
```

```
In [74]: prob = pd.DataFrame(predictedProbabilityScoresForEachClass)
prob.head()
```

```
Out[74]:
```

	0	1
0	0.185009	0.814991
1	0.254237	0.745763
2	0.947731	0.052269
3	0.927453	0.072547
4	0.199282	0.800718

```
In [75]: result = pd.DataFrame()
```

```
In [76]: result['patient_id']=test_data.index  
result['adherence'] = predictions_test
```

```
In [77]: result['prob_being_yes'] = predictedProbabilityScoresForEachClass[:,1]  
result['prob_being_no'] = predictedProbabilityScoresForEachClass[:,0]
```

```
In [78]: result.head()
```

```
Out[78]:
```

	patient_id	adherence	prob_being_yes	prob_being_no
0	1	1	0.814991	0.185009
1	2	1	0.745763	0.254237
2	3	0	0.052269	0.947731
3	4	0	0.072547	0.927453
4	5	1	0.800718	0.199282

```
In [79]: result['prob_score'] = result[["prob_being_yes", "prob_being_no"]].max(axis=1)
```

```
In [80]: result['adherence'] = result['adherence'].apply(lambda x: 'Yes' if x == 1 else 'No')
```

```
In [81]: result.head()
```

```
Out[81]:
```

	patient_id	adherence	prob_being_yes	prob_being_no	prob_score
0	1	Yes	0.814991	0.185009	0.814991
1	2	Yes	0.745763	0.254237	0.745763
2	3	No	0.052269	0.947731	0.947731
3	4	No	0.072547	0.927453	0.927453
4	5	Yes	0.800718	0.199282	0.800718

**Dropping individual class probability column as in the given format we need only the predicted class probability**

```
In [82]: result.drop(['prob_being_yes','prob_being_no'], axis=1, inplace=True)
```

```
In [83]: result.head()
```

```
In [83]: result.head()
```

```
Out[83]:
```

	patient_id	adherence	prob_score
0	1	Yes	0.814991
1	2	Yes	0.745763
2	3	No	0.947731
3	4	No	0.927453
4	5	Yes	0.800718

```
In [84]: result = result.set_index('patient_id')
```

```
In [85]: result.head()
```

```
Out[85]:
```

	adherence	prob_score
patient_id		
1	Yes	0.814991
2	Yes	0.745763
3	No	0.947731
4	No	0.927453
5	Yes	0.800718

```
In [86]: result.to_csv('result.csv')
```

## Download Links:

<https://ashutoshtripathi.com/reference-python-notebook-for-important-ml-algorithms-for-the-beginners/>

## Thank You

For any help and support in Data Science, Machine Learning and Artificial Intelligence,

Please connect with me on

Website: <https://enetwork.ai/> or <https://enetwork.ai/contact-us>

Blog: <https://ashutoshtripathi.com/>

LinkedIn: <https://www.linkedin.com/in/ashutoshtripathi1/>

Instagram: [https://www.instagram.com/ashutosh\\_ai/](https://www.instagram.com/ashutosh_ai/)

Medium Articles: <https://medium.com/@ashutosh.optimistic>